Operation count of a Measurement Equation
M.A. Brentjens

November 11, 2005

## Contents

1 Introduction ..... 3
2 Notation ..... 3
3 Measurement Equation ..... 4
4 Solving a timeslot ..... 4
4.1 Model patch visibilities ..... 4
4.2 At the start of a solver iteration ..... 5
4.3 Derivative evaluation ..... 5
5 Computational costs ..... 5
5.1 Elementary operations ..... 5
5.2 Partial Fourier kernel $k_{i p s r t}$ ..... 6
5.3 UVW coordinates ..... 6
5.4 Computing a patch visibility ..... 6
5.5 Corrupting a patch ..... 7
5.6 Adding all patches ..... 7
6 High level cost equation ..... 7
7 Software ..... 8
7.1 Coding convention ..... 8
7.2 Hardware class ..... 8
7.3 InstructionStats class ..... 8
7.4 Helper functions ..... 8
8 Examples ..... 9
8.1 WSRT observation 3C 343 ..... 10
8.2 Full LOFAR, 10 kHz channels ..... 11
8.3 Full LOFAR, 200 channels ..... 12
8.4 Full LOFAR, 10 patches ..... 13
8.5 LOFAR core, 32 stations, 10 seconds, 200 subbands ..... 14
8.6 LOFAR core, 32 stations, 100 patches ..... 15
9 Comparison with MeqTree system ..... 16
10 A bit more realistic ..... 18
11 How to reduce cost ..... 19
12 Conclusions ..... 19

Table 1: List of symbols

| Symbol | Explanation |
| :--- | :--- |
| $x$ | Scalar |
| $z^{*}$ | Complex conjugate of $z$ |
| $\mathbf{x}$ | Vector |
| $\mathbf{X}$ | $2 \times 2$ matrix |
| $\mathbf{a} \cdot \mathbf{b}$ | Vector inner product between a and $\mathbf{b}$ |
| $\mathrm{V}_{i j p r t}$ | Uncorrupted patch visibility |
| $\mathrm{J}_{i p t}$ | Station / sky patch Jones matrix |
| $\mathbf{l}_{p s}$ | "lmn" coordinate vector of source $s$ of patch $p$ |
| $\mathbf{u}_{i t}$ | "uvw" coordinate of station $i$ at timeslot $t$ |
| $c$ | Speed of light in vacuum |
| $\mathrm{E}_{p s r}$ | Normalized coherency matrix of source $s$ of patch $p$ at fre- |
| $\tilde{\mathrm{V}}_{i j t r}$ | quency channel $r$ |
|  | Fully corrupted model visibility on baseline between sta- |
| $t$ | tions $i$ and $j$ |
| $r$ | timeslot index (0-based) |
| $R$ | frequency channel index (0-based) |
| $N$ | number of frequency channels |
| $B$ | number of stations |
| $P$ | number of baselines |
| $P$ | number of patches |
| $F$ | number of beams with $R$ channels each |
| $S_{p}$ | number of sources in patch $p$ |
| $S$ | total number of sources |
| $\nu_{r}$ | central frequency of channel $r$ |
| $\delta \nu$ | channel frequency increment |

## 1 Introduction

This memo is very much work in progress and has not yet been verified by anyone besides the author.

ASTRON has been working on implementing a self calibration system for LOFAR since $2001 / 2002$. To my knowledge, no one has ever estimated the minimum amount of operations needed to evaluate a Measurement Equation in a more-or-less realistic scenario. In this memo I analyze the number of operations needed for the evaluation of a particular ME (predict) for the purpose of solving for instrumental or ionospheric corrections.

## 2 Notation

The notation adhered to in this memo is summarized in Tab. 1.

## 3 Measurement Equation

At a given timeslot $t$ and frequency channel $r$, the visibility on a baseline between station $i$ and station $j$ of a collection of patches $p$ and point sources $p s$ is given by

$$
\begin{equation*}
\tilde{\mathrm{V}}_{i j r t}=\sum_{p=1}^{P} \mathrm{~J}_{i p t} \mathrm{~V}_{i j p r t} \mathrm{~J}_{j p t}^{\dagger} \tag{1}
\end{equation*}
$$

Equation (1) is a very minimal direction dependent measurement equation. The frequency independent Jones matrix J defaults to the unit matrix. The model patch visibility matrix

$$
\begin{equation*}
\mathrm{V}_{i j p r t}=\sum_{s=1}^{S_{p}} k_{i p s r t} \mathrm{E}_{p s r} k_{j p s r t}^{*} \tag{2}
\end{equation*}
$$

where the normalized source coherency matrix

$$
\mathrm{E}_{p s r}=\frac{1}{n_{p s}}\left(\begin{array}{cc}
<x x>_{p s r} & <x y>_{p s r}  \tag{3}\\
<y x>_{p s r} & <y y>_{p s r}
\end{array}\right)
$$

The partial Fourier kernel $k_{i p s r t}$ is a complex scalar given by

$$
\begin{equation*}
k_{i p s r t}=\mathrm{e}^{-2 \pi \mathrm{i}\left(u_{i t} l_{p s}+v_{i t} m_{p s}+w_{i t}\left(n_{p s}-1\right)\right) \nu_{r} / c} \tag{4}
\end{equation*}
$$

## 4 Solving a timeslot

In this memo we assume that only two coefficients per $\mathrm{J}_{i p t}$ are solvable. For example the diagonal phases, diagonal amplitudes, off-diagonal amplitudes, or off-diagonal phases.

When solving for instrumental calibration coefficients, there are three ways of evaluating the measurement equation. First, one needs to compute the model patch visibilities. Then, at the start of each iteration, one should multiply all patch visibilities with the updated sets of Jones matrices. The last phase in every iteration is the computation of the derivatives of all visibilities to the solvable coefficients of the Jones matrices. The costs of adding equations to the solver and computing the matrix solution are not modelled in this memo.

The following sub sections treat these "modes" in more detail.

### 4.1 Model patch visibilities

If one attempts to solve for phases or amplitudes of J , the model patch visibility, is invariant. It must therefore be computed only once at the beginning of the solution of a timeslot.

If the channels $r$ are equidistant, the partial Fourier kernels $k_{i p s r t}$ can be computed very efficiently:

$$
\begin{align*}
k_{i p s r t} & =\mathrm{e}^{-2 \pi \mathrm{i}\left(u_{i t} l_{p s}+v_{i t} m_{p s}+w_{i t}\left(n_{p s}-1\right)\right)\left(\nu_{0}+r \delta \nu\right) / c}  \tag{5}\\
k_{i p s r t} & =k_{i p s t}^{\nu_{0}} \times\left(k_{i p s t}^{\delta \nu}\right)^{r} \tag{6}
\end{align*}
$$

where

$$
\begin{equation*}
k_{i p s t}^{\nu_{0}}=\mathrm{e}^{-2 \pi \mathrm{i}\left(u_{i t} l_{p s}+v_{i t} m_{p s}+w_{i t}\left(n_{p s}-1\right)\right) \nu_{0} / c} \tag{7}
\end{equation*}
$$

and

$$
\begin{equation*}
k_{i p s t}^{\delta \nu}=\mathrm{e}^{-2 \pi \mathrm{i}\left(u_{i t} l_{p s}+v_{i t} m_{p s}+w_{i t}\left(n_{p s}-1\right)\right) \delta \nu / c} . \tag{8}
\end{equation*}
$$

After this, one needs one complex scalar-scalar multiplication and one complex scalar-matrix multiplication in order to compute the model visibiity of one source. Then one requires $P\left(S_{p}-1\right)$ matrix additions in order to compute $\mathrm{V}_{i j p r t}$ for all patches. This must be repeated for all $B$ baselines and $R$ channels. It is assumed that all baselines are used, so

$$
\begin{equation*}
B=\frac{N(N-1)}{2} . \tag{9}
\end{equation*}
$$

That is nevertheless seldom the case in reality. Usually, the shorter baselines are not used in the calibration because of difficulties in modelling extended sources.

### 4.2 At the start of a solver iteration

The first step of a solver iteration is a full prediction, given the current set of Jones matrices. Given the invariant $\mathrm{V}_{i j p r t}$, one needs to multiply all patches with their current Jones matrices and add them to form a visibility point. This involves two complex matrix-matrix multiplications per patch, and $P-1$ complex matrix additions. This must be repeated for all $B$ baselines and $R$ channels.

### 4.3 Derivative evaluation

For all visibility points, one needs to evaluate the derivative of the measurement equation with respect to all solvable parameters. We approximate derivatives by forward differencing. That means that for each derivative, one needs a partial evaluation of the measurement equation. Because each coefficient is specific to one patch only, one must multiply this particular patch with the affected Jones matrices and add all patches to obtain the perturbed visibility. The matrix multiplication must be done for all channels and all $N-1$ baselines involving the perturbed Jones matrix.

## 5 Computational costs

### 5.1 Elementary operations

All basic math involves complex numbers. The operations needed for basic algebra are listed in the following table:

| Operation | multiply | add |
| :--- | ---: | ---: |
| compl. scalar-scalar mul | 4 | 2 |
| compl. scalar-matrix mul | 16 | 8 |
| compl. matrix-matrix mul | 32 | 24 |
| compl. matrix-matrix add | 0 | 8 |
| real matrix-3-vector mul | 9 | 6 |

### 5.2 Partial Fourier kernel $k_{i p s r t}$

Computing $l$, $m$, and $n-1$ is done per source and is done once per calibration. Those costs are therefore negligible compared to all other operations. I consider them non existent.

The cost of evaluating $k_{i p s r t}$ for all $R$ channels is:
mul: $7+4(R-1)+1$
add: $3+2(R-1)$
sin: 2
cos: 2
The " +1 " in the mul is the multiplication by $1 / c$.
The partial Fourier kernel must be computed for all station-source combinations. That is, that is, $N P S_{p}$ times for each timeslot.

### 5.3 UVW coordinates

A very rough estimate, assuming $63 \times 3$ rotation matrices (three "nominal" rotations, three correction rotations). One needs to compute, for each matrix, a sine and cosine. Some time calculations (Taylor series?) are also required. Let's estimate that one needs 20 terms, yielding 20 more multiplications and 20 more additions for the date/time computations.

The total cost for one UVW coordinate computation would then be estimated as
mul: 74
add: 56
sin: 6
cos: 6
Because the uvw coordinates are computed only once for all stations ( $N$ times per timeslot, independent of baseline and frequency) their computing cost is almost irrelevant.

### 5.4 Computing a patch visibility

Given that $\mathrm{E}_{p s r}, k_{i p s r t}$, and $k_{j p s r t}$, are already available, one needs $S_{p}$ complex scalar-scalar multiplications and $S_{p}$ complex scalar-matrix multiplications to predict all source visibilities. Combining these source visibilities into a patch visibility costs an additional $S_{p}-1$ complex matrix additions. The cost for one patch visibility is then
mul: $20 S_{p}$
add: $18 S_{p}-8$
This needs to be done $P R B$ times per timeslot. It is not necessary to repeat this step for each iteration as the sky model is held constant.

### 5.5 Corrupting a patch

Corrupting a patch visibility for one baseline, one channel costs two complex matrix-matrix multiplications.
mul: 64
add: 48
This needs to be done $P R B$ times at the start of each solver iteration.
For each $J_{i p t}$ parameter derivative one has to re-corrupt all patch visibilities that involve this parameter. That means at most $(N-1) R$ times per parameter per timeslot.

If there are two solvable parameters per patch per station in one solution: $2(N-1) P \times B R$ times per timeslot per solver iteration for phases and $2 N P \times B R$ for amplitudes.

### 5.6 Adding all patches

Adding all patches costs $P-1$ complex matrix-matrix additions.
mul: 0
add: $8 P-8$
This must be repeated for all baselines and channels at the start of every iteration. That is, $B R$ times. In the derivative calculations, the patches must be added once for each derivative. That means $(2(N-1) P+1) B R$ times in total per timeslot for phase solutions and $(2 N P+1) B R$ for amplitudes. Therefore, this component to the total computation time scales with the square of the number of patches! Fortunately, it is a relatively cheap operation. It nevertheless is the limiting factor when the number of patches is larger than approximately 20.

## 6 High level cost equation

The total cost for computing the patch visibilities is equal to the cost of the $u v w$, $k_{i p s r t}$, source visibilities $\mathrm{V}_{i j p s r t}$, and the combination of all source visibilities into patch visibilities $\mathrm{V}_{i j p r t}$.

At the start of each iteration, one must corrupt $P$ patches and combine them all into the unperturbed visibility $\mathrm{V}_{i j r t}$. Then one must corrupt $N-1$ patches and combine them for all $\approx 2 N P$ solvables.

A few relevant lines from the Python script are

```
insSkyModel = insUVW + insStatSourceDFT + insSourcePatchVis +\
    insAddSources
insFullEveryIteration = insCorruptPatch*cPatches + insAddPatches
insDerivatives = insCorruptPatch*((cStations-1.0)/cBaselines)\
    + insAddPatches
ins = (insSkyModel + (insFullEveryIteration +\
    (insDerivatives*cSolvablesPerTimeslot)\
    )*cIterations)*cBeams
```

Currently, the MeqTree solver typically converges in 4 iterations. Because only 2 out of 8 coefficients are solved for in one iteration, one needs to do this 4 times. That is, in the end one needs of order 16 iterations. If there are multiple beams with $R$ channels, but in independent directions, then the entire cost must be multiplied by the number of beams.

## $7 \quad$ Software

I have written a small utility in Python that does the actual operations counting. I first introduce the coding convention. Then I proceed to explain which classes exist and how they interact.

### 7.1 Coding convention

I use apps Hungarian (Charles Simonyi, Microsoft). Prefixes used in the script:
c Count of something
s String
ins InstructionStats instance
hw Hardware instance
cst Cost of an operation in CPU clock cycles
rf Frequency in Hz

### 7.2 Hardware class

The Hardware class holds a description of a CPU. It is used to help converting instruction counts to clock cycles and execution times in seconds. The actual conversion is done in InstructionStats. SFromHw ()

### 7.3 InstructionStats class

The InstructionStats class helps with instruction arithmetic. It holds counts for the multiply, add, sine and cosine instructions. One can add two InstructionStats instances, or multiply an InstructionStat instance with a real number.

The InstructionStats. SFromHw (hw) method converts instruction counts to clock cycles, given a Hardware instance hw.

### 7.4 Helper functions

There are several "constructor" functions that construct InstructionStats instances for useful cases, such as a scalar-scalar product, a matrix-matrix product, a smart, and a stupid implementation of $k_{i p s r t}$.

## 8 Examples

This section contains a number of example outputs of the predict_runtime.py script. The first table always gives the configuration that was used in the computations. The second table gives the operation in the first column, the number of times the operation is called in the second column, and the number of clock cycles in the third column. For the clock cycle computation, the following CPU model was used.:
name: AMD 1900+
clock: 1.6 GHz
mul: 1 cycle
add: 1 cycle
sin: 20 cycles
cos: 20 cycles
The results that are quoted are the following:
SkyModel: the model patch visibilities $\mathrm{V}_{i j p r t}$
CorruptPatches: $\mathrm{J}_{i p t} \mathrm{~V}_{i j p r t} \mathrm{~J}_{j p t}^{\dagger}$
AddPatches: adding all patches to form $\mathrm{V}_{i j r t}$
The amount of CPU seconds needed to compute all predict costs for one timeslot is given below the count table. Finallyi, the minimum amount of TFlops needed to do all ME evaluations and derivative computations in real time is quoted. This number assumes $100 \%$ CPU efficiency adn does not include adding equations to the solver and performing the solver step itself.

### 8.1 WSRT observation 3C 343

| Settings: |  |  |  |
| :---: | :---: | :---: | :---: |
| Paramete |  | Value |  |
| Stations |  | 14.0 |  |
| Channels | per beam | 16.0 |  |
| Patches |  | 2.0 |  |
| Sources p | er patch | 1.0 |  |
| Sources in | total | 2.0 |  |
| Solvables | per timeslot | 52.0 |  |
| Solver Ite | rations | 4 |  |
| Beams |  | 1 |  |
| Integratio | n time | 30.0 |  |
| Results: |  |  |  |
| Component | Instructions | Cycles |  |
| SkyModel |  |  |  |
| Multiply | 61180 | 61180 |  |
| Add | 30828 | 30828 |  |
| Cosine | 140 | 2800 |  |
| Sine | 140 | 2800 |  |
| CorruptPatch |  |  |  |
| Multiply | 3514367 | 3514367 |  |
| Add | 2635775 | 2635775 |  |
| Cosine | 0 | 0 |  |
| Sine | 0 | 0 |  |
| AddPatches |  |  |  |
| Multiply | 0 | 0 |  |
| Add | 2469376 | 2469376 |  |
| Cosine | 0 | 0 |  |
| Sine | 0 | 0 | $+$ |
| Multiply | 3575548 | 3575548 |  |
| Add | 5135980 | 5135980 |  |
| Cosine | 140 | 2800 |  |
| Sine | 140 | 2800 |  |

Time@AMD 1900+: 0.00545 seconds

### 8.2 Full LOFAR, 10 kHz channels

Settings:

| Parameter | Value |
| :--- | :--- |
| Stations | 77.0 |
| Channels per beam | 3200.0 |
| Patches | 20.0 |
| Sources per patch | 500.0 |
| Sources in total | 10000.0 |
| Solvables per timeslot | 3040.0 |
| Solver Iterations | 16 |
| Beams | 1 |
| Integration time | 1.0 |

Results:

| Results: |  |  |
| :--- | ---: | ---: |
| Component | Instructions | Cycles |
| SkyModel |  |  |
| Multiply | 1882499085698 | 1882499085698 |
| Add | 1688806662312 | 1688806662312 |
| Cosine | 1540462 | 30809240 |
| Sine | 1540462 | 30809240 |
| CorruptPatch |  |  |
| Multiply | 948830208000 | 948830208000 |
| Add | 711622656000 | 711622656000 |
| Cosine | 0 | 0 |
| Sine | 0 | 0 |
| AddPatches |  |  |
| Multiply | 0 | 0 |
| Add | 09247530598400 | 69247530598400 |
| Cosine | 0 | 0 |
| Sine | 0 | 0 |
|  |  |  |
| Multiply | 2831329293698 | 2831329293698 |
| Add | 71647959916712 | 71647959916712 |
| Cosine | 1540462 | 30809240 |
| Sine | 1540462 | 30809240 |

Time@AMD 1900+: 46549.5942681 seconds Real-time calibration requires at least 74.5 TFlops.

### 8.3 Full LOFAR, 200 channels

| Settings: |  |
| :--- | :--- |
| Parameter | Value |
| Stations | 77.0 |
| Channels per beam | 200.0 |
| Patches | 20.0 |
| Sources per patch | 500.0 |
| Sources in total | 10000.0 |
| Solvables per timeslot | 3040.0 |
| Solver Iterations | 16 |
| Beams | 1 |
| Integration time | 1.0 |

Results:

| Component | Instructions | Cycles |
| :--- | ---: | ---: |
| SkyModel |  |  |
| Multiply | 117659085698 | 117659085698 |
| Add | 105551142312 | 105551142312 |
| Cosine | 1540462 | 30809240 |
| Sine | 1540462 | 30809240 |
| CorruptPatch |  |  |
| Multiply | 59301888000 | 59301888000 |
| Add | 44476416000 | 44476416000 |
| Cosine | 0 | 0 |
| Sine | 0 | 0 |
| AddPatches |  |  |
| Multiply | 0 | 0 |
| Add | 0 | 0 |
| Cosine | 4327970662400 | 4327970662400 |
| Sine | 0 | 0 |
|  |  | 0 |
| Multiply | 176960973698 | 176960973698 |
| Add | 4477998220712 | 4477998220712 |
| Cosine | 1540462 | 30809240 |
| Sine | 1540462 | 30809240 |

Time@AMD 1900+: 2909 seconds
Real-time calibration requires at least 4.7 TFlops.

### 8.4 Full LOFAR, 10 patches

| Settings: |  |
| :--- | :--- |
| Parameter | Value |
| Stations | 77.0 |
| Channels per beam | 200.0 |
| Patches | 10.0 |
| Sources per patch | 1000.0 |
| Sources in total | 10000.0 |
| Solvables per timeslot | 1520.0 |
| Solver Iterations | 16 |
| Beams | 1 |
| Integration time | 1.0 |

Results:

| Component | Instructions | Cycles |
| :--- | ---: | ---: |
| SkyModel |  |  |
| Multiply | 117659085698 | 117659085698 |
| Add | 105597958312 | 105597958312 |
| Cosine | 1540462 | 30809240 |
| Sine | 1540462 | 30809240 |
| CorruptPatch |  |  |
| Multiply | 29650944000 | 29650944000 |
| Add | 22238208000 | 22238208000 |
| Cosine | 0 | 0 |
| Sine | 0 | 0 |
| AddPatches |  |  |
| Multiply | 0 | 0 |
| Add | 0 | 0 |
| Cosine | 1025382758400 | 1025382758400 |
| Sine | 0 | 0 |
|  |  | 0 |
| Multiply | 147310029698 | 147310029698 |
| Add | 1153218924712 | 1153218924712 |
| Cosine | 1540462 | 30809240 |
| Sine | 1540462 | 30809240 |

Time@AMD 1900+: 813 seconds
Real-time calibration requires at least 1.3 TFlops.

### 8.5 LOFAR core, 32 stations, 10 seconds, 200 subbands



Time@AMD 1900+: 5322 seconds
Real-time calibration requires at least 0.9 TFlops.

### 8.6 LOFAR core, 32 stations, 100 patches

|  | Settings: |  |
| :---: | :---: | :---: |
|  | Parameter | Value |
|  | Stations | 32.0 |
|  | Channels per beam | 200.0 |
|  | Patches | 100.0 |
|  | Sources per patch | 100.0 |
|  | Sources in total | 10000.0 |
|  | Solvables per timeslot | 6200.0 |
|  | Solver Iterations | 16 |
|  | Beams | 24 |
|  | Integration time | 10.0 |
| Results: |  |  |
| Component | Instructions | Cycles |
| SkyModel |  |  |
| Multiply | 482334776832 | 482334776832 |
| Add | 429719083008 | 429719083008 |
| Cosine | 15364608 | 307292160 |
| Sine | 15364608 | 307292160 |
| CorruptPatch |  |  |
| Multiply | 1188495360000 | 1188495360000 |
| Add | 891371520000 | 891371520000 |
| Cosine | 0 | 0 |
| Sine | 0 | 0 |
| AddPatches |  |  |
| Multiply | 0 | 0 |
| Add | 187081054617600 | 187081054617600 |
| Cosine | 0 | 0 |
| Sine | 0 | 0 |
|  |  |  |
| Multiply | 1670830136832 | 1670830136832 |
| Add | 188402145220608 | 188402145220608 |
| Cosine | 15364608 | 307292160 |
| Sine | 15364608 | 307292160 |

Time@AMD 1900+: 118796 seconds
Real-time calibration requires at least 19.0 TFlops.


Figure 1: Measured execution time per timeslot of MeqTree system.

## 9 Comparison with MeqTree system

In order to compare the runtime prediction to the actual execution time of the MeqTree system, I did a timing experiment. The execution time per timeslot solution was measured as a function of the number of timeslots that were solved simultaneously in the same solver instance.

This approach enabled a decomposition in several distinct execution time contributions. First there is a fixed overhead involved with, e.g., function calls. This overhead is constant per solve domain. Its contribution per timeslot is proportional to $1 /($ number of timeslots). Then there is the inversion of the covariance matrix in the solver. Its cost is proportional to the third power of the number of solvable parameters. As the number of solvables is linearly proportional to the number of timeslots in a solution domain, the contribution of the matrix inverse is proportional to the number of timeslots cubed. Last but not least there is the fixed amount of work that needs to be done for every single timeslot. It consists of

- The actual predict, including any inefficiencies
- Adding normal equations to the solver
- I/O and data handling
- Something else?

In this experiment we can not discriminate between these contributions.

The measurements are plotted in Fig. 1. The dashed line is the best least squares fit of Eqn. (10) to the data.

$$
\begin{equation*}
t=\left(\frac{a}{x}+b x^{3}+c\right) n_{\text {iter }} \tag{10}
\end{equation*}
$$

The best fit coefficients are $a=0.60 \pm 0.03, b=2.13 \pm 0.05 \times 10^{-5}, c=$ $0.059 \pm 0.004$. For our comparison, coefficient $c$ is the most interesting. Referring back to Sec. 8.1, one sees that at least 0.00545 s are required for the predict on an AMD1900+, which coincidentally is precisely the machine on which this experiment was run.

For four iterations, we see that the constant amount of time spent per timeslot is $0.059 \times 4=0.24 \mathrm{~s}$. This is 43 times slower than the minimum computed in Sec. 8.1. At the moment we do not know what the contribution of adding the equations is.

We are able to estimate the amount of time spent in disk I/O. The dataset is 1.3 GB in size. If we assume that the effective disk I/O rate is of the order of 20 MB per second, and we consider that the data are also corrected and written back to disk, then a grossly pessimistic estimate of the disk I/O contribution is: 0.09 seconds per timeslot. (read AND write full $1.3 \mathrm{~GB}, 1440$ timeslots in total).

Subtracting this disk-I/O leaves 0.15 seconds for computing. If we furthermore assume that one can attain effectively $50 \%$ of the peak performance of the CPU, then we are down to 0.07 seconds. Still almost a factor of 13 more than the estimate in Sec. 8.1.

When analyzing the tree profiling data we consistently see that the solver consumes most of the CPU cycles. We snevertheless still need to investigate whether that is caused by adding the normal equations or by other, constant-per-timeslot overhead in the solver.

Using coefficient $b$ of Eqn. (10), we can estimate the actual cost of the matrix inversion inside the Aips++ Levenberg-Marquardt solver. For $n$ solvable parameters, the matrix inversion costs: $1.14 \times 10^{-9} n^{3} \mathrm{~s}$ on an AMD $1900+\mathrm{CPU}$. That corresponds to about $1.82 \times n^{3} \mathrm{CPU}$ cycles.

If there are 77 stations and only 2 coefficients are solved for per matrix inversion, then one matrix inversion would cost about 6.6 million CPU cycles. On average, four iterations are needed to converge. Assuming 20 patches, and considering that we need to to this four times per matrix in order to determine the other coefficients, and multiplying the result by 200 channels/subbands, we find that in this case 420 Gcycles are needed to do all the matrix inversions for one timeslot. If one cycle corresponds to one Flop, 420 GFlops should be reserved for matrix inversions in the solver. This does not include adding the normal equations, which could possibly be far more expensive than the relatively small matrix inversions.

The makeNorm equation we use:

```
template <class U, class V, class W>
void LSQFit::makeNorm(uInt nIndex, const W &cEqIndex,
    const V &cEq, const U &weight,
    const U &obs,
    Bool doNorm, Bool doKnown) {
    if (doNorm) {
```

```
        for (uInt i=0; i<nIndex; ++i) {
            if (cEq[i] != 0) {
                Double *i2 = norm_p->row(cEqIndex[i]); //row pointer
                Double eq(cEq[i]);
                    eq *= weight;
                for (uInt i1=0; i1<nIndex; ++i1) {
                if (cEqIndex[i]<=cEqIndex[i1] && cEq[i1] != 0) {
                        i2[cEqIndex[i1]] += eq*Double(cEq[i1]); //equations
                };
                };
            };
        };
        state_p &= "TRIANGLE;
    };
    if (doKnown) {
        Double obswt = obs*weight;
        for (uInt i1=0; i1<nIndex; ++i1) {
            if (cEq[i1] != 0) {
                known_p[cEqIndex[i1]] += Double(cEq[i1])*obswt; //data vector
            };
        };
        error_p[NC] += 1; //cnt equations
        error_p[SUMWEIGHT] += weight; //sum weight
        error_p[SUMLL] += obs*obswt; //sum rms
    };
}
```


## 10 A bit more realistic

Evidently, the amount of computing power needed for the self calibration is astonishing. In order to obtain a more realistic number for the required amount of TFlops, one should take into account the cost of the makeNorm() function call, approximately $50 \%$ CPU efficiency, and the duty cycle for the given experiment.

An estimate for the cost of one makeNorm call is of the order of 50 cycles for the call itself, and 7 times the number of derivatives for a given datapoint squared. There are typically $4 P$ derivatives per data point, so we would have a cost of the order of $112 P^{2}$ cycles per data point. In total this needs to be done $B R$ times per iteration per beam. Therefore the contribution of makeNorm to a full lofar with 20 patches and 200 channels that would amount to 420 GFlops at $100 \%$ efficiency.

Therefore an estimate for the selfcal cost of LOFAR could be: (4.7 TFlops + 0.4 TFlops $($ makeNorm $)+0.4$ TFlops(solver matrix invert $)) \times 2($ efficiency $) \times 2($ anything we forgot business) $=11$ TFlops.

For the EOR a similar estimate would be: $19+$

## 11 How to reduce cost

Polynomials help in reducing the cost of evaluating the measurement equation. In fact, they make matters worse. Instead of "one" derivative per datapoint, one must compute " $N_{\text {coeff }}$ derivatives, effectively multiplying the required processing power by the degree of the polynomial plus one.

A coarser data grid is the easiest structural thing that helps considerably. It is already evident that it would be rediculous to calibrate on a 1 second, 10 kHz grid. Instead I propose to calibrate per subband in 1 s time intervals for the full LOFAR. PErhaps we even need to go to 10 s timeslots. In any case we must properly account for time- and bandwidth smearing in the presence of missing data points at a finer grid than the calibration grid.

Correction should nevertheless be performed on a finer grid. Here it might help to use some form of polynomial interpolation for accuracy. The correction step is probably much cheaper because it needs only be done once per beam per timeslot and not for all derivatives and iterations.

One can also reduce the number of baselines in the solution:

- no short baselines: get rid of Milkyway / large scale, hard to model stuff
- no long baselines: Little signal anyway because sources are resolved.


## 12 Conclusions

In terms of sheer FLOPS, calibrating the full LOFAR in real time is going to be extremely challenging. Chances are much better for the EOR experiment that only uses 32 stations and 10 seconds integration time in combination with a low dutycycle. All in all I think that ways can be found to do the calibration in almost real time for reasonable processing cost. We may however need to make the calibration cluster a bit more powerful than initially intended.

